



University
of Windsor

GENG-8900 MACHINE LEARNING

Instructor: Dr. Yasser Alginahi

K-Means Clustering

Harshkumar Zaveri

Aeshita Dhiman


Gurpreet Singh Saini

Date: 03-Nov-2023

Overview

- Types of Machine Learning
- Clustering
- What is K-means clustering
- Key Terminologies
- K means algorithm
- Evaluate quality of clusters
- Advantages and Disadvantages
- Implementation and code
- Applications
- References


Types of Machine Learning



Types of Machine Learning [<https://shorturl.at/pzNT9>]

Clustering

- Clustering is the process of arranging a group of objects in such a manner that the objects in the same group are more similar to each other than to the objects in any other group




Clustering [<https://shorturl.at/hpyDE>]

Clustering

- 5 essential clustering algorithms
 - K-Means
 - MeanShift
 - DBSCAN
 - Hierarchical clustering
 - BIRCH

What is K-mean clustering?

- K-means clustering is an unsupervised machine learning algorithm that partitions a dataset into K clusters
- Such models aim to assign similar data points into distinct clusters or groups



K-means clustering [<https://shorturl.at/kvPS6>]

Key Terminologies

- In ideal conditions for clustering to perform well the inter-cluster distance should be very high and the intra-cluster distance should be very low.
- Based on these two distances we can measure the performance of clustering techniques using the Dunn index.
- The Dunn Index is the ratio of the smallest distance between observations not in the same cluster to the largest intra-cluster distance.



Inter-cluster and Intra-cluster distance [<https://shorturl.at/fzBU2>]

Key Terminologies

- **Data Points:** These are individual data instances or observations that are grouped into clusters based on their similarities.
- **Centroids:** Centroids are the central points of each cluster. They represent the mean or average of the data points within that cluster.
- **Clusters:** Clusters are groups of data points that are similar to each other and dissimilar to data points in other clusters. Each data point belongs to one and only one cluster.
- **K-value:** K refers to the number of clusters you want to create in the dataset. It is a hyperparameter that you must choose before running the K-means algorithm.



Keywords [<https://shorturl.at/fzBU2>]


Key Terminologies

- This algorithm takes the data points and gives the predefined set of clusters. Consider dataset $D=\{x_1, x_2, \dots, x_n\}$ has n points with k clusters(S_1, S_2, \dots, S_k) and each cluster assigned with one centroid(C_1, C_2, \dots, C_k).
- Every point in the dataset is assigned to the set (S_i) depending on the nearest centroid to it.
- The centroid(C_i) is the mean point to all the other points in the given set of points (S_i). It is calculated as below:

$$C_i = \frac{1}{n} \sum_{X_j \in S_i} X_j$$

K-Means Algorithm

- Step 1: Identify number of clusters
- Step 2: Initialize Cluster Centroids
- Step 3: Assign Datapoints to Centroids
- Step 4: Re-evaluate cluster mean
- Step 5: Repeat steps 3 – 4 till convergence or a certain number of iterations




K-means Algorithm [<https://shorturl.at/ouMOW>]

K-Means Algorithm

Identify number of clusters

- Elbow Method
 - Find optimal number of clusters using Sum of Squared Distances (SSD)
- Methodology
 - Start with a defined range of K
 - Run the KMeans Algorithm on your dataset
 - Calculate SSD – Euclidean distance
 - The sum of the squared distance from centroid.
 - $d = (x_1 - c_1)^2 + (x_2 - c_2)^2$
 - Sum up the squares
 - Repeat for all clusters

$$\arg \min_{(c_1, c_2 \dots c_k)} = \sum_{i=1}^K \sum_{x \in S_i} ||x - c_i||^2$$



Elbow Method [\[https://shorturl.at/gIJM8\]](https://shorturl.at/gIJM8)

K-Means Algorithm

Assign initial Centroids

- At Random
 - Basic method
- KMeans ++
 - Smart centroid initialization
 - Initial centroid assigned randomly
 - Rest of the centroids based on maximum squared distance to spread out centroids

K-Means Iteration [<https://shorturl.at/jwWZ7>]



K-Means Algorithm

Assign data points to clusters

- Using distance measures like Euclidian distance to find the nearest centroid to the data point
- What Happens if same distance?
 - Choose any cluster at random

K-Means Iteration [<https://shorturl.at/jwWZ7>]

K-Means Algorithm

ITERATION 1

Assign data points to nearest
centroid to create clusters

K-Means Iteration [<https://shorturl.at/jwWZ7>]



K-Means Algorithm

Update centroids and clusters

- Recalculate the centroid for each cluster using :

$$\bullet \quad c_j = \frac{1}{|s_j|} \sum_{x_i \in s_j} x_i$$

- Reassign data points to the new centroids

K-Means Iteration [<https://shorturl.at/jwWZ7>]



K-Means Algorithm

Repeat data point assignment and centroid update until convergence

- Convergence :The state where there is a very small update in the new centroids over the old centroids

K-Means Iteration [<https://shorturl.at/jwWZ7>]



Evaluate quality of Cluster

- **Inertia**

- It tells us how far the points within a cluster are
- the sum of distances of all the points within a cluster from the centroid of that cluster
- Euclidean distance is often used as the distance metric
- We calculate this for all the clusters; the final inertial value is the sum of all these distances
- the distance between them should be as low as possible

- **Dunn Index**

- Dunn index is the ratio of the minimum of inter-cluster distances and maximum of intra-cluster distances.
- The more the value of the Dunn index, the better the clusters will be

$$\text{Dunn Index} = \frac{\min(\text{Inter cluster distance})}{\max(\text{Intra cluster distance})}$$




Clusters are far apart



Clusters are compact

Advantages


- Easy to implement
- Result can be Interpreted seamlessly
- Produce consistent results
- Can adapt to new examples
- The biggest advantage:
UNSUPERVISED LEARNING




K-mean clustering graph [<https://tinyurl.com/2s3zsdrx>]

Disadvantages

- Choosing K manually
- Clustering outliers
- Different random seeds can cause the algorithm to converge to different local minima instead of the global minimum
- It assumes that clusters are spherical and have similar variance



(A): Undesirable clusters



(B): Ideal clusters

Effect of outlier on clustering [<https://tinyurl.com/c29dyzc4>]

Implementing K-Mean on a real dataset

- The [dataset](#) is from an automobile company, which have divided there existing customer into four groups and want to predict the group for new customer.
- The dataset have following columns: Unique ID, Gender, Marital status, Age, Graduated, profession, Work experience, Credit score, Family size and category group
- What are we going to do? Remove the category that's already defined and cluster it by ourself.

Code!!!

```
Import pandas as pd  
from sklearn.cluster import Kmeans
```

```
df = pd.read_csv("Train.csv")  
df = df.drop(["Segmentation", "ID"], axis="columns")  
df=df.dropna()  
df = df.reset_index()  
df = df.drop("index", axis="columns")  
df_new = pd.get_dummies(df)  
df_kmeans = pd.get_dummies(df, drop_first=True)  
kmeans_model = KMeans(n_clusters=3)  
clusters = kmeans_model.fit_predict(df_kmeans)  
df_kmeans.insert(df_kmeans.columns.get_loc("Age"), "Cluster", clusters)  
df_kmeans.sample(5, random_state=44)
```

K-mean Implementation and code [<http://bit.ly/49kADke>]



Implementation

Read the dataset

```
df = pd.read_csv("Train.csv")
df.sample(5, random_state=44)
```

	ID	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	Spending_Score	Family_Size	Var_1	Segmentation
1627	466681	Male	No	19	No	Healthcare	2.0	Low	4.0	Cat_6	D
805	462758	Male	No	18	No	Healthcare	0.0	Low	3.0	Cat_6	D
3725	467212	Female	No	28	Yes	Doctor	1.0	Low	3.0	Cat_6	B
6581	460636	Male	Yes	45	Yes	Artist	2.0	Average	3.0	Cat_3	D
7462	459157	Male	No	37	Yes	Engineer	1.0	Low	1.0	Cat_6	D



Implementation

Remove columns which we do not need

```
df = df.drop(["Segmentation", "ID"], axis="columns")
df.sample(5, random_state =40)
```

	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	Spending_Score	Family_Size	Var_1
3757	Female	No	38	No	Homemaker	13.0	Low	NaN	Cat_6
4692	Female	Yes	45	Yes	Artist	1.0	High	5.0	Cat_6
5629	Male	Yes	80	Yes	Lawyer	NaN	Low	1.0	Cat_6
341	Male	Yes	51	No	Entertainment	1.0	Average	4.0	Cat_7
6150	Female	Yes	50	Yes	Artist	1.0	Average	3.0	Cat_6

Implementation

Detail information about the dataset

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8068 entries, 0 to 8067
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Gender            8068 non-null    object  
 1   Ever_Married      7928 non-null    object  
 2   Age                8068 non-null    int64  
 3   Graduated          7990 non-null    object  
 4   Profession         7944 non-null    object  
 5   Work_Experience    7239 non-null    float64 
 6   Spending_Score     8068 non-null    object  
 7   Family_Size         7733 non-null    float64 
 8   Var_1               7992 non-null    object  
dtypes: float64(2), int64(1), object(6)
memory usage: 378.3+ KB
```



Implementation

Data Cleaning

```
df=df.dropna()  
df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 6665 entries, 0 to 8067  
Data columns (total 9 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   Gender            6665 non-null    object    
 1   Ever_Married     6665 non-null    object    
 2   Age               6665 non-null    int64     
 3   Graduated         6665 non-null    object    
 4   Profession       6665 non-null    object    
 5   Work_Experience   6665 non-null    float64   
 6   Spending_Score    6665 non-null    object    
 7   Family_Size       6665 non-null    float64   
 8   Var_1              6665 non-null    object    
dtypes: float64(2), int64(1), object(6)  
memory usage: 364.5+ KB
```



Implementation

Dataset after Data cleaning and preprocessing

```
df = df.reset_index()
df = df.drop("index", axis="columns")
df.head()
```

	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	Spending_Score	Family_Size	Var_1
0	Male	No	22	No	Healthcare	1.0	Low	4.0	Cat_4
1	Female	Yes	67	Yes	Engineer	1.0	Low	1.0	Cat_6
2	Male	Yes	67	Yes	Lawyer	0.0	High	2.0	Cat_6
3	Male	Yes	56	No	Artist	0.0	Average	2.0	Cat_6
4	Male	No	32	Yes	Healthcare	1.0	Low	3.0	Cat_6

Implementation

Formatting the data for K-Mean clustering

```
df_new = pd.get_dummies(df)
df_new.head()
```

	Age	Work_Experience	Family_Size	Gender_Female	Gender_Male	Ever_Married_No	Ever_Married_Yes	Graduated_No	Graduated_Yes	Pro
0	22	1.0	4.0	0	1	1	0	1	0	0
1	67	1.0	1.0	1	0	0	1	0	0	1
2	67	0.0	2.0	0	1	0	1	0	0	1
3	56	0.0	2.0	0	1	0	1	1	1	0
4	32	1.0	3.0	0	1	1	0	0	0	1

Implementation

Formatting the data for K-Mean clustering

```
df_kmeans = pd.get_dummies(df, drop_first=True)  
df_kmeans.head()
```

	Age	Work_Experience	Family_Size	Gender_Male	Ever_Married_Yes	Graduated_Yes	Profession_Doctor	Profession_Nurse
0	22	1.0	4.0	1	0	0	0	0
1	67	1.0	1.0	0	1	1	1	0
2	67	0.0	2.0	1	1	1	1	0
3	56	0.0	2.0	1	1	0	0	0
4	32	1.0	3.0	1	0	1	0	0

Implementation

Create K-mean clustering model

```
kmeans_model = KMeans(n_clusters=3)
clusters = kmeans_model.fit_predict(df_kmeans)
df_kmeans.insert(df_kmeans.columns.get_loc("Age"), "Cluster", clusters)
df_kmeans.sample(5, random_state=44)
```

	Cluster	Age	Work_Experience	Family_Size	Gender_Male	Ever_Married_Yes	Graduated_Yes	Profession_Docor	Profession_Engineer	Profession_Entertainment	.
6434	0	36	13.0	1.0	0	0	1	0	0	0	0
1190	0	18	2.0	4.0	1	0	0	0	0	0	0
1298	2	40	9.0	2.0	0	0	1	0	0	0	0
6591	0	35	9.0	2.0	1	1	1	0	0	0	1
6054	1	77	0.0	2.0	0	1	1	0	0	0	0

Applications

- Image Compression
- Image segmentation
- Evaluating performance – Academic, work or credit



Image compression using K-mean [<http://tiny.cc/52gdvz>]




Image segmentation using K mean [<http://tiny.cc/52gdvz>]

References

- What is K-means Clustering [<https://rb.gy/eqh727>]
- K-means Algorithm [<https://rb.gy/anbipp>]
- Elbow Method [<https://rb.gy/jdjn28>]
- K-mean Implementation and code [<http://bit.ly/49kADke>]
- Advantages and Disadvantages [<https://bit.ly/40qKkda>]

